



Systematic Evaluation of Signal-to-Image Transformation Pipelines for Vision Transformer-based Sleep Stage Classification

Janith R.R.H. Ramanayakage and Chandima N.P.G. Arachchige

Department of Statistics, Faculty of Science, University of Colombo, Sri Lanka

Why automate sleep staging?



Sleep disorders (insomnia, apnoea, narcolepsy) are widespread



Gold standard: manual scoring of polysomnography (PSG)



Slow and subjective: hours per night; experts disagree



Deep learning: CNN/RNN → Vision Transformers (ViTs)

2–4 h

to score one full-night recording

82.6%

inter-scorer agreement, experienced experts

50–70 M

US adults living with a sleep disorder

Research gap and objectives

The gap

- Prior ViT work uses spectrograms or raw signals only
- GASF, GADF and MTF never compared for ViTs
- Best sampling rate before imaging unknown
- Pre-training value on PSG images unclear

Four experiments

- 1 Image transformation: GASF, GADF, MTF, spectrogram
- 2 Sampling rate: 512 / 256 / 128 / 64 Hz, FFT vs DWT
- 3 Pre-trained vs from-scratch ViT, balanced vs imbalanced
- 4 Seven architectures incl. ResNet-50 baseline

Dataset: CAP Sleep Database

80

recordings (6 healthy,
74 patients)

80,667

30-second epochs


16

channels at 512 Hz

6

stages: Wake, S1–S4,
REM

Subject-wise 5-fold CV: no subject in both
train and test



Proposed Pipeline




Figure 1. Overall five-stage pipeline

Four Signal-to-Image transformations



GASF

$$\cos(\varphi_i + \varphi_j)$$

Summed angles;
preserves global
temporal correlation



GADF

$$\sin(\varphi_i - \varphi_j)$$

Difference of angles;
less globally coherent



MTF

**Markov transition
probs**

32 bins; quantisation
loss, slow (5 s/epoch)



Spectrogram

STFT, Hamming 256


Fixed time–frequency
resolution



16 images → 4×4 grid → 224×224×3 → ViT-B/16

GASF performs best

EXPERIMENT ● ○ ○ ○



78.5%

GASF accuracy (macro-F1
0.773)


5.0 s

MTF time per epoch vs 0.2 s for
GASF

⚡ GASF $\approx$ GADF; both clearly beat MTF and spectrogram

DWT at 128 Hz is optimal

EXPERIMENT ●●○○



80.3%

DWT @ 128 Hz (macro-F1
0.791)


+1.8 pp

over original 512 Hz signals

⚡ DWT beats FFT at every reduced rate; peak at 128 Hz

Pre-training adds 11+ points

EXPERIMENT ●●●○



+11.4 pp

balanced: 86.2% vs 74.8%


+12.2 pp

imbalanced: 80.5% vs 68.3%

⚡ Balancing (DCGAN, training folds only) adds ≈ 5.7 pp

ViT-B/16 beats other models

EXPERIMENT ●●●●



86.2%

ViT-B/16 (macro-F1 0.853)


+10.0 pp

over ResNet-50 CNN baseline



Transformers outperform the CNN; PVT-Small is the efficient option

Class-wise results and limitations



Limitations

- S1 (6.5% of data) is weakest
- Single database; no cross-dataset test
- R&K labels, not AASM
- No formal significance tests
- 4×4 channel layout not ablated

Conclusions and future work

Best pipeline

-  **GASF** image encoding
-  **DWT** resampling to 128 Hz
-  **Fine-tuned** ImageNet ViT-B/16
-  **86.2%** accuracy, 0.853 macro-F1

Future work

-  Multimodal fusion
-  Lightweight ViTs for edge use
-  Cross-dataset validation



Thank you

Questions?

janithramanayake523@gmail.com